

WENQI CHEN

chenwenqi@std.uestc.edu.cn

(+86) 18876290529

No. 4, Section 2, North Jianshe Road, Chengdu, Sichuan, P.R.China, 610054

EDUCATION

The Hong Kong University of Science and Technology (HKUST)

Expected Fall 2026

Department of Computer Science and Engineering (CSE)

Incoming Ph.D. Student

University of Electronic Science and Technology of China (UESTC)

August 2022 - Present

School of Information and Software Engineering

Bachelor's degree in progress

Key Courses: *Calculus, Linear Algebra, Numerical Analysis, Computer Simulation, Operating Systems, Data Structures and Design, Computer Organization Principles*

Current GPA: 3.91/4.00 | **Weight Average:** 87.1/100

RESEARCH EXPERIENCE

HKGAI, The Hong Kong University of Science and Technology

June 2024 - January 2025

Research Collaborator

Topic: Post-training Value Alignment for Regional Sovereign LLM (HKGAI-V1)

- Spearheaded the construction of region-specific value datasets and optimized data mixtures using the *Align-Anything* framework, ensuring compliance with Hong Kong's safety standards.
- Implemented fine-grained LLF to upgrade reward modeling from scalar scoring to sequence generation, effectively mitigating "safety refusals" and improving utility for administrative tasks.
- The model is deployed in **30+ government departments** and featured in the *2024 Policy Address*. Presented technical innovations to NVIDIA CEO Jensen Huang and HKSAR officials.

PAIR Lab, Institute for Artificial Intelligence, Peking University

June 2024 - Present

Research Intern

Topic: AI Alignment in All-modality Models & Framework Development

- Co-developed *Align-Anything*, a widely-used cross-modal alignment framework with **4.6k+ GitHub stars**. Contributed to Learning from Language Feedback (LLF) algorithms and paper writing.
- Constructed a 200k all-modality human preference dataset with fine-grained language feedback. The dataset has achieved **46,800+ downloads**, becoming a standard resource in the community.

Topic: Mitigating Deceptive Alignment in Large Language Models

- Developed *CoT Monitor+*, a mechanism embedding self-monitoring signals into Chain-of-Thought to intercept deceptive reasoning in real-time.
- Constructed *DeceptionBench* (180 scenarios) to evaluate 5 categories of deceptive behaviors. Experiments demonstrated a **43.8% reduction** in deception while preserving task accuracy.

- Modeled mitigation as constrained RL to bridge the evaluation-deployment gap. The interdisciplinary analysis on human perception of deception is under review at *Nature Human Behaviour*.

PUBLICATION

- **Wenqi Chen**, Ziyang Zhang, Bing Wang, Lin Liu, Hengheng Zhang, Zhengsu Chen. "Learn from Your Mistakes: Tree-like Self-Play for Secure Code LLMS." *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, 2026.
- Isabella Duan, Xudong Pan, Yawen Duan, Adam Gleave, Ranjie Duan, Chaochao Lu, Naying Hu, Sören Mindermann, Dongrui Liu, Jie Fu, Xudong Guo, Chen Zheng, **Wenqi Chen**, Jianfeng Cao, Geng Hong, Brian Tse, Xia Hu, Min Yang, Yang Zhang, Xiaojian Li, Peng Xu, Tianxing He, Jiarun Dai, Yinpeng Dong. "Position: Preparing for AI Systems That Deceive Developers." *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, 2026.
- Li, Chao, Ziwei Deng, Chenxing Lin, **Wenqi Chen**, Yongquan Fu, Weiquan Liu, Chenglu Wen, Cheng Wang, Siqi Shen. "Dof: A diffusion factorization framework for offline multi-agent reinforcement learning." ICLR 2025.
- Ji, Jiaming, Jiayi Zhou, Hantao Lou, Boyuan Chen, Donghai Hong, Xuyao Wang, **Wenqi Chen** ... & Yaodong Yang. "Align anything: Training all-modality models to follow instructions with language feedback." arXiv preprint arXiv:2412.15838 (2024).
- Ji, Jiaming*, **Wenqi Chen***, Kaile Wang, Donghai Hong, Sitong Fang, Boyuan Chen, Jiayi Zhou, Juntao Dai, Sirui Han, Yike Guo, Yaodong Yang. "Mitigating deceptive alignment via self-monitoring." arXiv preprint arXiv:2505.18807 (2025).
- Zhou, Jiayi*, Jiaming Ji*, Boyuan Chen, Jiapeng Sun, **Wenqi Chen**, Donghai Hong, Sirui Han, Yike Guo, Yaodong Yang. "Generative RLHF-V: Learning Principles from Multi-modal Human Preference." NeurIPS 2025.
- Chen, Boyuan*, Donghai Hong*, Jiaming Ji*, Jiacheng Zheng, Bowen Dong, Jiayi Zhou, Kaile Wang, Juntao Dai, Xuyao Wang, **Wenqi Chen**, ..., Yike Guo, Yaodong Yang. "InterMT: Multi-Turn Interleaved Preference Alignment with Human Feedback." NeurIPS 2025 **Spotlight**.

PRACTICAL SERVICE

Chain Era Student Studio, UESTC

October 2023 - Present

Vice Chairman & Head of AI Division

Managed competitions, research, and daily affairs of the studio.

Led recruitment and knowledge sharing sessions for the machine learning division.

HONORS & AWARDS

UESTC Excellent Student Scholarship

(2022-2023-2024)

Second Prize, Sichuan Province College Student Computer Design Competition

(2024)

National Third Prize, China Collegiate Computing Contest - Network Technology Challenge. (2025)